

Intelligent Medical Report Analysis System

Project Domain: DL / GenAI / NLP
DSAI LAB Group 2



Indian Institute Of Technology, Madras

DSAI LAB Group 2

21F2001310

Shivendra Patel

21f2001310@ds.study.iitm.ac.in

Signature: S.P.

23F1001316

Samta Ranka

23f1001316@ds.study.iitm.ac.in

Signature: S.R.

22F3003195

Rajat Srivastava

22f3003195@ds.study.iitm.ac.in

Signature: R.S.

22F1000120

Ritwik Trivedi

22f1000120@ds.study.iitm.ac.in

Signature: R.T.W.

22F3002277

Bryan David Robinson

22f3002277@ds.study.iitm.ac.in

Signature: B.D.R.

Final Contribution Summary

Intelligent Medical Report Analysis System — Group No. 2

TEAM DISTRIBUTION 5 Key Specialists		SYSTEM PIPELINE OCR → NER → LLM Summary		LEADERSHIP STRUCTURE Rotational across M1–M6		SHARED OWNERSHIP Report, PPT & QA Reviews		
Ritwik Trivedi 22f1000120 PRIMARY FOCUS LLM Fine-Tuning, Chatbot Implementation & LLM Documentation KEY DELIVERABLES - LLM Architecture Selection - Fine-Tuning & Evaluating LLM - LLM Chatbot Feature Implementation & Review - Deployment on Hugging Face		Rajat Srivastava 22f3003195 Milestone 5 Lead PRIMARY FOCUS OCR Integration, End-to-End Pipeline Setup, Model Performance & PPT KEY DELIVERABLES - OCR Preprocessing & Workflow Validation - End-to-End Integration (OCR + Extraction + LLM) - Inference Output & Error Analysis Reports		Samta Ranka 23f1001316 Milestone 4&6 Lead PRIMARY FOCUS Extraction Model, Mimic-III Preprocessing & Overall Documentation KEY DELIVERABLES - Dataset Preprocessing, Annotation - ClinicalBERT Fine-Tuning - OCR Preprocessing & Workflow Validation - Front-End UI Development - PPT		Shivendra Patel 21f2001310 Milestone 1-3 Lead PRIMARY FOCUS Extraction Model, Mimic-III Preprocessing & Overall Documentation KEY DELIVERABLES - Dataset Preprocessing, Annotation - ClinicalBERT Fine-Tuning - OCR Preprocessing & Workflow Validation - Overall M1-M6 Documentation		Bryan David Robinson 22f3002277 PRIMARY FOCUS LLM Fine-Tuning, Chatbot Implementation & LLM Documentation KEY DELIVERABLES - LLM Architecture Selection & LLM Dataset Preprocessing - Hyperparameter Optimization & Training Logs - Final Evaluation Artifacts - LLM Chatbot Feature Implementation & Review
Rotational Leadership: Shivendra (M1-3) → Samta (M4,6) → Rajat (M5) Shared Ownership: Report Compilation, PPT Preparation & Quality Review across all Milestones								

Problem Definition & Target Objectives

Clinical Problem & Jargon Barrier

Medical lab reports use complex medical terms and shorthand, making them difficult to understand and increasing patient anxiety, confusion, and risk of misinterpretation.

Generic AI Hallucination & Cloud PHI Risk

LLMs can produce unstructured text, hallucinate critical values, and raise HIPAA/GDPR privacy concerns when PHI is sent to public cloud APIs.

Scope Boundaries & Market Benchmarks

In-Scope Functionality

Process lab reports (PDF/Image), extract entities (test, value, unit), flag against reference bounds, plain-English summary, local offline execution.

Out-of-Scope Boundaries

No clinical diagnosis/prescriptions, no medical imaging (X-Ray/CT/MRI), no handwritten notes. Proof-of-concept for educational comprehension.

Objective	Target	Status Achieved
Extraction Accuracy	Entity F1 ≥ 0.85	0.9991 (EXCEEDED)
Abnormality Flagging	Deterministic ranges	ZERO HALLUCINATION
Readability Level	FKGL below 8	GRADE 8.4 (MEAN)
Output System	Dual Output	DASHBOARD + NARRATIVE

Platform	Key Strengths	Core Limitations
Med-PaLM	High medical QA accuracy	Proprietary & Closed-Source
MedGemma	Multimodal health understanding & open weights	Requires clinical fine-tuning for lab extraction
Glass Health	Differential diagnosis	Clinician focus, no patient UI
UpToDate / AMBOSS	Trusted medical databases	Manual lookup, requires clinical background
Our System	Automated, 100% private	GPU dependency

Problem Solution – ClinicalBERT Extraction and BioMistral-7B for Zero Hallucinations and generating Patient-Friendly Insights



- **ClinicalBERT** extracts lab entities and converts them into structured JSON, keeping extraction separate from text generation.
- A rule-based module compares values with reference ranges to classify results as HIGH , LOW or NORMAL .
- **BioMistral-7B** only converts the verified structured results into patient-friendly explanations, **reducing numerical hallucination**.

Exploratory Data Phase

Started with Kaggle MIMIC-III sample (100 patient records, 76,074 LABEVENTS rows) + MTSamples (2,416 lab transcriptions).

Full PhysioNet Access & Scaling

Expanded to full PhysioNet access (1.65M records). Ablation studies proved a 20,000 stratified report subset delivered identical F1 accuracy while reducing training time from >10 hrs to 30 mins.

MTSAMPLES EXTRACTION

1,630 Unique

- Total MTSamples Reports **4,999 total**
- Lab-Related Keywords **582 keywords**
- Lab Reports Identified **2416 lab tests related**
- Deduplicated Reports **1,630 unique**

FULL MIMIC-III MINING

Total 1.6 Million Reports

- Total LABEVENTS Rows **27.85M events**
- Valid Numerical Results **24.93M valid**
- Patients with Lab Data **46,250 patients**
- D_LABITEMS Rows **753 (589 unique)**
- Structured Lab Entities **231,988 entities**

FINAL CORPUS (70/15/15)

21,628 Reports

- MIMIC Synthetic Share **19,998 (92.5%)**
- MTSamples Share **1,630 (7.5%)**
- Training Set (70%) **15,138 reports**
- Validation Set (15%) **3,245 reports**
- Test Set (15%) **3,245 reports**
- Patient Leakage Audit **Passed (0)**

TRAIN / VALIDATION / TEST SPLIT (70/15/15)



- Train: 15,138 (70%)
- Val: 3,245 (15%)
- Test: 3,245 (15%)

DATA SET SOURCE DISTRIBUTION



END TO END DATA PROCESSING PIPELINE



NER Annotation & Tokenization for ClinicalBERT

2.4M Tokens Annotated

Dictionary preparation, BIO tag classification scheme, and token distribution plots

1. Dictionary & Annotation

589

MIMIC-III LABELS

1,859

SEARCHABLE VARIATIONS

Rule-based annotation engine

Mimic Labels + regex patterns auto-labelled medical entities across the corpus - replacing slow manual labelling.

Tag-Set Refinement

9 IOB tags → pruned to 5: {B-TEST, I-TEST, B-VALUE, B-UNIT, O}.

Fixed class imbalance & unit/range confusion via loss re-weighting.

Switched model selection to Entity F1 (not token accuracy).

2. NER Labels & Concrete Example

Label	Meaning / Entity Category
B-TEST	Beginning of test name
I-TEST	Continuation of test name
B-VALUE	Numerical laboratory result
B-UNIT	Laboratory measurement unit
O	Other / non-entity token

Anion

B-TEST

Gap

I-TEST

=

O


22

B-VALUE


mEq/L

B-UNIT


Misclassification-Only Heatmap




Annotated Entity Distribution



BIO Label Distribution



System Architecture Design & Pipeline



ClinicalBERT — NER Extractor Model Specs

Architecture: 110M params · 12 layers · 768 hidden · 12 heads

Pretraining Domain: Pretrained on MIMIC-III clinical notes

Task Head: Token classification head (5 BIO tags)


Hyperparameters: lr 2e-5 · 10 epochs · Best @ Epoch 9 (val loss 0.0016)

Extraction F1 Score: 99.91% (entity-level)

ClinicalBERT Training Metrics 10 Epochs Evaluation


Epoch	Train Loss	Val Loss	Accuracy	Entity F1
Epoch 1	0.01584	0.00912	0.9979	0.9979
Epoch 3	0.00384	0.00318	0.9993	0.9993
Epoch 5	0.00115	0.00185	0.9996	0.9996
Epoch 9 BEST	0.00054	0.00166	0.9997	0.9991

ClinicalBERT Performance by Report Length



Dataset Size	Train Time	Entity F1	Decision
1.6M	>10 h	99%	Slow
21628	0.5 h	99.91%	Selected
5,000	0.1 h	98.2%	Slight drop

ClinicalBERT Training & Validation Loss Trend



Real-World Images/PDF OCR Challenges & Fixes

Challenge - Zero Lab Tests Identified using Fine-Tuned ClinicalBert on Raw Text extracted from Lab Tests Images/PDF using OCR

Fixes



1. Rotation Detection & Correction

Canny + Hough Lines detect and correct image rotation.



2. Image Preprocessing

Upscaling, CLAHE, and sharpening enhance image quality.



3. PaddleOCR Extraction

Extracts text and bounding boxes from the report.



4. Bounding-Box Line Sorting

Reconstructs OCR tokens into correctly ordered text lines.




5. Value-Line Filtering

Keeps relevant laboratory lines containing values and test keywords.



6. Lexicon Normalization

Corrects OCR misspellings using medical test-name patterns.



DARSS Composite Metric & Success

COMPOSITE EVALUATION FORMULA

$DARSS = (0.45 \times Text_{Det}) + (0.35 \times Val_{Acc}) + (0.20 \times Unit_{Acc}) + Bonus$

85%

FULL SUCCESS
DARSS ≥ 70%

0%

PARTIAL
DARSS 50–69%

15%

UNSUCCESSFUL
DARSS < 50%

Overall Error Rate

Text Detection Error 15.3%

15.3%

Value Extraction Error 14.8%

14.8%

Unit Extraction Error 22.2%

22.2%

Overall Model Performance

Precision

98.2%

Recall

83.8%

F1-Score

89.3%

Value Accuracy

85.2%


Unit Accuracy

77.8%

Format Evaluation Breakdown

Image Format Type	Count	Det %	Val %	Unit %
Screenshot	3	100.0%	100.0%	86.1%
Rotated Image	4	88.2%	88.2%	72.7%
PDF Report	2	81.8%	90.0%	70.0%
Human Captured	3	77.0%	77.0%	87.7%
Low-Quality Image	3	76.9%	87.0%	69.8%
Standard Lab	5	69.0%	69.0%	61.7%

System Architecture Design & Pipeline



BioMistral-7B QLoRA Explainer Model Specs & Metrics

- **Quantization:** 4-bit (NF4 + double quant) · bfloat16
- **LoRA Config:** Rank r=8, α =16 · 20.97M params (0.2888%)
- **Selection:** Outperformed Mistral, Phi-3, Gemma-3, Qwen-3

Metric	Value	Interpretation / Note
Eval NLL Loss	0.6906	Lowest loss among 8 trials
Perplexity	1.99	eNLL · High token certainty
Eval Token Acc.	80.75%	Agreement with reference text
Train-Eval Gap	0.0587	0.6906 - 0.6319 · Strong general.

Lowest Zero-Shot FKGL + 100% Numeric Fidelity

8-Trial QLoRA Exploration Matrix Selected Candidate: Trial 5

Trial #	Rank (r)	Alpha (α)	LR	Train Loss	Eval Loss	Token Acc.	Status / Finding
Trial 0	16	16	1.05e-4	0.8808	0.7697	78.54%	Moderate perf.
Trial 1	32	32	7.79e-5	0.8047	0.7427	79.35%	Slower conv.
Trial 2	32	128	2.18e-3	0.8324	0.7847	78.52%	Underfit (low LR)
Trial 3	64	256	2.36e-4	0.5425	0.7403	80.33%	Overfit (high cap)
Trial 4	32	64	2.75e-3	0.9491	0.8515	76.80%	Highest eval loss
Trial 5 (Best)	8	16	3.31e-4	0.6319	0.6906	80.75%	Optimal balance
Trial 6	16	64	2.56e-4	0.5168	0.7131	81.07%	Slight overfit
Trial 7	32	128	2.18e-5	0.8331	0.7851	78.55%	Underfit (low LR)

100%

Numeric Fidelity

Every value, unit & status reproduced exactly

0.52

ROUGE-L

Strong content overlap with reference summaries

0.40

BLEU

Respectable n-gram precision for open generation


8.4

FKGL (target ≤ 8)


Readability is the known shortfall to improve



1 Image



2 Dashboard



3 Explanations

Interactive Explanation

- Hemoglobin (15.2 g/dL, Normal):** Hemoglobin carries oxygen in your blood. This level is normal.
- Hematocrit (44.1 %, Normal):** Hematocrit measures the amount of blood that is made up of red blood cells. This result is normal.
- Platelet Count (280.0 x10³/μL, Normal):** Platelets help your blood clot to stop bleeding. This count is normal.
- MCV (89.0 fL, Normal):** MCV measures the size of your red blood cells. This value is normal.
- Neutrophils (62.0 %, Normal):** Neutrophils are white blood cells that fight infection. This percentage is normal.
- Glucose (98.0 mg/dL, Normal):** Glucose is a type of sugar in your blood. This level is normal.
- Sodium (141.0 mmol/L, Normal):** Sodium helps control the amount of water in your body. This level is normal.
- Potassium (4.3 mmol/L, Normal):** Potassium is a mineral that helps your body work properly. This level is normal.
- Chloride (105.0 mmol/L, Normal):** Chloride helps keep the balance of fluids and acids in your body. This level is normal.
- Bicarbonate (26.0 mmol/L, Normal):** Bicarbonate helps keep the balance of acids and bases in your body. This level is normal.
- Bun (16.0 mg/dL, Normal):** Bun is a type of nitrogen that helps your body work properly. This level is normal.
- Creatinine (1.05 mg/dL, Normal):** Creatinine is a waste product that your kidneys remove from your blood. This level is normal.
- Cholesterol, Total (215.0 mg/dL, High):** Cholesterol is a type of fat that can build up in your blood. A high level may suggest a problem with your heart or blood vessels, so it is worth checking

4 Q&A

Analytical Overview

AI Medical Assistant
Online & ready

Explain my high Cholesterol, Total

Your Cholesterol, Total is 215 mg/dL. That is higher than the normal range. The normal range is <200 mg/dL.

Note: reference applies to: desirable
If this concerns you, bring it up with your doctor.

Explain my high Cholesterol, Total

Explain my high Cholesterol, LDL, Calculated

Ask a question about your report...

AI guidance is not a diagnosis. Discuss with your doctor.

 Deployment Infrastructure

- **GPU-enabled Deployment:** Inference backend deployed on **Hugging Face Spaces** for high-throughput BioMistral processing.
- **Local Execution Support:** Full end-to-end pipeline capable of running locally on suitable GPU hardware.
- **Self-Hosted Models:** ClinicalBERT & BioMistral run on dedicated endpoints without third-party LLM dependencies.

 Privacy & Data Flow

- **Zero Persistent Storage:** Medical reports processed entirely in-memory and never stored permanently.
- **Transient Lifespan:** Data retained strictly for the duration of OCR, NER extraction, and explanation.
- **Isolated Transmission:** Data routes exclusively to the dedicated inference backend - never sent to public APIs.

 Medical Safety & Guardrails

- **Non-Diagnostic Design:** System provides analytical explanations only; explicitly defers medical diagnosis to clinicians.
- **Medication Guardrails:** Clear warnings against altering or stopping prescribed medication without a physician.
- **Emergency Deferral:** High-risk or emergency indicators trigger immediate prompts for professional clinical care.

 Architecture Comparison Matrix

Approach	Entity F1	Hallucination Risk	Data Privacy	Decision Rationale
Pretained ClinicalBERT)	8.20%	None	High	Not Selected
Heuristic Regex	24.73%	None	High	Not Selected
★ Our Hybrid Pipeline	99.91%	ZERO HALLUCINATION	TRANSIENT PROCESSING	SELECTED PIPELINE



THANK YOU

